

# \_\_lab5

October 2, 2024

0.0.1

-

0.1

№ 5

cc

0.1.1

:

:

: -01-21

0.2

2024

0.2.1

№25

25

Ozone Level Detection Data Set

: eighthr.data

: <http://archive.ics.uci.edu/ml/datasets/Ozone+Level+Detection>

: class ( No 74)

– (SelectKBest)

:

•

•

•

(degree=2)

- ROC-

1.

UCI,

.

```
[1]: import numpy as np
```

```
[2]: import pandas as pd
```

```
[3]: import matplotlib.pyplot as plt
```

```
[4]: import warnings
warnings.filterwarnings("ignore")
```

```
[5]: # DataFrame
my_data = pd.read_csv( 'eighthr.data', header=None )
my_data
```

```
[5]:
```

	0	1	2	3	4	5	6	7	8	9	...	64	\
0	1/1/1998	0.8	1.8	2.4	2.1	2	2.1	1.5	1.7	1.9	...	0.15	
1	1/2/1998	2.8	3.2	3.3	2.7	3.3	3.2	2.9	2.8	3.1	...	0.48	
2	1/3/1998	2.9	2.8	2.6	2.1	2.2	2.5	2.5	2.7	2.2	...	0.6	
3	1/4/1998	4.7	3.8	3.7	3.8	2.9	3.1	2.8	2.5	2.4	...	0.49	
4	1/5/1998	2.6	2.1	1.6	1.4	0.9	1.5	1.2	1.4	1.3	...	?	
...	...	...	...	...	...	...	...	...	...	...	...	...	
2529	12/27/2004	0.3	0.4	0.5	0.5	0.2	0.3	0.4	0.4	1.3	...	0.07	
2530	12/28/2004	1	1.4	1.1	1.7	1.5	1.7	1.8	1.5	2.1	...	0.04	
2531	12/29/2004	0.8	0.8	1.2	0.9	0.4	0.6	0.8	1.1	1.5	...	0.06	
2532	12/30/2004	1.3	0.9	1.5	1.2	1.6	1.8	1.1	1	1.9	...	0.25	
2533	12/31/2004	1.5	1.3	1.8	1.4	1.2	1.7	1.6	1.4	1.6	...	0.54	
...	...	...	...	...	...	...	...	...	...	...	...	...	
0	10.67	-1.56	5795	-12.1	17.9	10330	-55	0	0.0				
1	8.39	3.84	5805	14.05	29	10275	-55	0	0.0				
2	6.94	9.8	5790	17.9	41.3	10235	-40	0	0.0				
3	8.73	10.54	5775	31.15	51.7	10195	-40	2.08	0.0				
4	?	?	?	?	?	?	?	0.58	0.0				
...	...	...	...	...	...	...	...	...	...	...	...	...	
2529	7.93	-4.41	5800	-25.6	21.8	10295	65	0	0.0				
2530	5.95	-1.14	5845	-19.4	19.1	10310	15	0	0.0				
2531	7.8	-0.64	5845	-9.6	35.2	10275	-35	0	0.0				
2532	7.72	-0.89	5845	-19.6	34.2	10245	-30	0.05	0.0				
2533	13.07	9.15	5820	1.95	39.35	10220	-25	0	0.0				

[2534 rows x 74 columns]

```
[6]: my_data = my_data.rename(columns={73: "class"})
my_data
```

```
[6]:
```

	0	1	2	3	4	5	6	7	8	9	...	64	\
0	1/1/1998	0.8	1.8	2.4	2.1	2	2.1	1.5	1.7	1.9	...	0.15	
1	1/2/1998	2.8	3.2	3.3	2.7	3.3	3.2	2.9	2.8	3.1	...	0.48	

2	1/3/1998	2.9	2.8	2.6	2.1	2.2	2.5	2.5	2.7	2.2	...	0.6
3	1/4/1998	4.7	3.8	3.7	3.8	2.9	3.1	2.8	2.5	2.4	...	0.49
4	1/5/1998	2.6	2.1	1.6	1.4	0.9	1.5	1.2	1.4	1.3	...	?
...	...	...	...	...	...	...	...	...	...	...	...	...
2529	12/27/2004	0.3	0.4	0.5	0.5	0.2	0.3	0.4	0.4	1.3	...	0.07
2530	12/28/2004	1	1.4	1.1	1.7	1.5	1.7	1.8	1.5	2.1	...	0.04
2531	12/29/2004	0.8	0.8	1.2	0.9	0.4	0.6	0.8	1.1	1.5	...	0.06
2532	12/30/2004	1.3	0.9	1.5	1.2	1.6	1.8	1.1	1	1.9	...	0.25
2533	12/31/2004	1.5	1.3	1.8	1.4	1.2	1.7	1.6	1.4	1.6	...	0.54

	65	66	67	68	69	70	71	72	class
0	10.67	-1.56	5795	-12.1	17.9	10330	-55	0	0.0
1	8.39	3.84	5805	14.05	29	10275	-55	0	0.0
2	6.94	9.8	5790	17.9	41.3	10235	-40	0	0.0
3	8.73	10.54	5775	31.15	51.7	10195	-40	2.08	0.0
4	?	?	?	?	?	?	?	0.58	0.0
...	...	...	...	...	...	...	...	...	...
2529	7.93	-4.41	5800	-25.6	21.8	10295	65	0	0.0
2530	5.95	-1.14	5845	-19.4	19.1	10310	15	0	0.0
2531	7.8	-0.64	5845	-9.6	35.2	10275	-35	0	0.0
2532	7.72	-0.89	5845	-19.6	34.2	10245	-30	0.05	0.0
2533	13.07	9.15	5820	1.95	39.35	10220	-25	0	0.0

[2534 rows x 74 columns]

2.

,  
,  
,  
( ),

```
[7]: #
my_data['class'].isnull().sum(axis=0)
```

[7]: 0

```
[8]: #
len(my_data['class'].unique())
```

[8]: 2

3.

- ,  
( ) ,

```
[9]: #
my_data.dtypes
```

```
[9]: 0      object
      1      object
      2      object
      3      object
      4      object
      ...
      69     object
      70     object
      71     object
      72     object
      class   float64
      Length: 74, dtype: object
```

```
[10]: #      NaN
      my_data[:73] = my_data[:73].replace('?', np.nan)
```

```
[11]: my_data = my_data.drop(my_data.columns[0], axis=1)
      my_data.dtypes
```

```
[11]: 1      object
      2      object
      3      object
      4      object
      5      object
      ...
      69     object
      70     object
      71     object
      72     object
      class   float64
      Length: 73, dtype: object
```

```
[12]: for col in my_data.columns[:73]:
      my_data[col] = pd.to_numeric(my_data[col], errors='coerce')

      #      NaN
      for col in my_data.columns[:73]:
          my_data[col] = my_data.groupby('class')[col].transform(lambda x: x.fillna(x.
      ↪mean()))
```

```
[13]: #      :      object
      #      float
      my_data.dtypes
```

```
[13]: 1      float64
      2      float64
      3      float64
```

```

4         float64
5         float64
...
69        float64
70        float64
71        float64
72        float64
class     float64
Length: 73, dtype: object

```

```

[14]: #
      #
      my_data.isnull().sum(axis=0)

```

```

[14]: 1         0
      2         0
      3         0
      4         0
      5         0
      ..
      69        0
      70        0
      71        0
      72        0
      class     0
      Length: 73, dtype: int64

```

```

4.
[15]: y = my_data['class']
      X = my_data.drop(columns='class')
      X.shape, y.shape

```

```

[15]: ((2534, 72), (2534,))

```

```

[16]: from sklearn import preprocessing
      X1 = X.to_numpy()
      X_scaled = preprocessing.scale(X1)
      new_data = pd.DataFrame(X_scaled)
      new_data

```

```

[16]:
      0         1         2         3         4         5         6  \
0   -0.706100  0.175988  0.729019  0.502307  0.420123  0.502524 -0.128527
1    0.966740  1.350171  1.500432  1.031026  1.572305  1.499570  1.151769
2    1.050382  1.014690  0.900444  0.502307  0.597382  0.865086  0.785970
3    2.555938  1.853393  1.843283  2.000345  1.217787  1.408930  1.060319
4    0.799456  0.427599  0.043319 -0.114532 -0.554800 -0.041319 -0.402876
...      ...      ...      ...      ...      ...      ...

```

2529	-1.124310	-0.998196	-0.899520	-0.907611	-1.175206	-1.129006	-1.134474
2530	-0.538816	-0.159493	-0.385244	0.149827	-0.023024	0.139962	0.145822
2531	-0.706100	-0.662715	-0.299532	-0.555132	-0.997947	-0.857084	-0.768675
2532	-0.287890	-0.578844	-0.042394	-0.290772	0.065605	0.230602	-0.494326
2533	-0.120606	-0.243363	0.214744	-0.114532	-0.288912	0.139962	-0.037077

	7	8	9	...	62	63	64 \
0	-0.319853	-0.575293	-0.479037	...	-1.314553	-0.637391	0.086034
1	0.686889	0.500175	0.476524	...	-1.050979	0.721427	-0.163600
2	0.595367	-0.306426	-0.305299	...	-1.419982	1.215542	-0.322358
3	0.412323	-0.127181	0.215917	...	-1.657199	0.762603	-0.126374
4	-0.594419	-1.113027	-1.260860	...	-0.037286	0.026451	0.044503
...	...	...	...	...	...	...	...
2529	-1.509638	-1.113027	-0.565906	...	-0.497475	-0.966801	-0.213964
2530	-0.502897	-0.396048	-0.392168	...	-0.392045	-1.090330	-0.430752
2531	-0.868984	-0.933782	-1.173991	...	-0.339331	-1.007977	-0.228198
2532	-0.960506	-0.575293	-0.739645	...	-0.075757	-0.225628	-0.236957
2533	-0.594419	-0.844160	0.129048	...	-0.365688	0.968484	0.348807

	65	66	67	68	69	70	71
0	-0.340272	-0.307405	-1.122041	-1.780072	3.223762	-1.582211	-0.282455
1	0.426052	-0.178248	0.175691	-0.766119	2.154203	-1.582211	-0.282455
2	1.271847	-0.371983	0.366752	0.357450	1.376342	-1.149766	-0.282455
3	1.376861	-0.565718	1.024303	1.307460	0.598481	-1.149766	1.296679
4	0.037422	-0.037417	-0.006215	-0.010517	0.025822	0.001992	0.157881
...	...	...	...	...	...	...	...
2529	-0.744721	-0.242826	-1.791998	-1.423818	2.543133	1.877344	-0.282455
2530	-0.280669	0.338379	-1.484314	-1.670455	2.834831	0.435863	-0.282455
2531	-0.209713	0.338379	-0.997975	-0.199767	2.154203	-1.005618	-0.282455
2532	-0.245191	0.338379	-1.494239	-0.291114	1.570807	-0.861470	-0.244495
2533	1.179604	0.015487	-0.424789	0.179323	1.084644	-0.717322	-0.282455

[2534 rows x 72 columns]

5. `10`, `25`: `(SelectKBest)`

```
[17]: from sklearn.feature_selection import SelectKBest, f_classif
valid_columns = [col for col in new_data.columns if new_data[col].nunique() >=
    ↪ 10 and col != 'class']

#
selector = SelectKBest(score_func=f_classif, k=2)
X_new = selector.fit_transform(X, y)

#
selected_indices = selector.get_support(indices=True)
```

```
#
selected_features = X.columns[selected_indices]
print("          :", selected_features)
```

```
: Index([42, 43], dtype='object')
```

```
[18]: #          ,          10
len(new_data[42].unique())
```

```
[18]: 339
```

```
[19]: #          ,          10
len(new_data[43].unique())
```

```
[19]: 331
```

```
[20]: #          2
data2 = new_data[[42,43]]
```

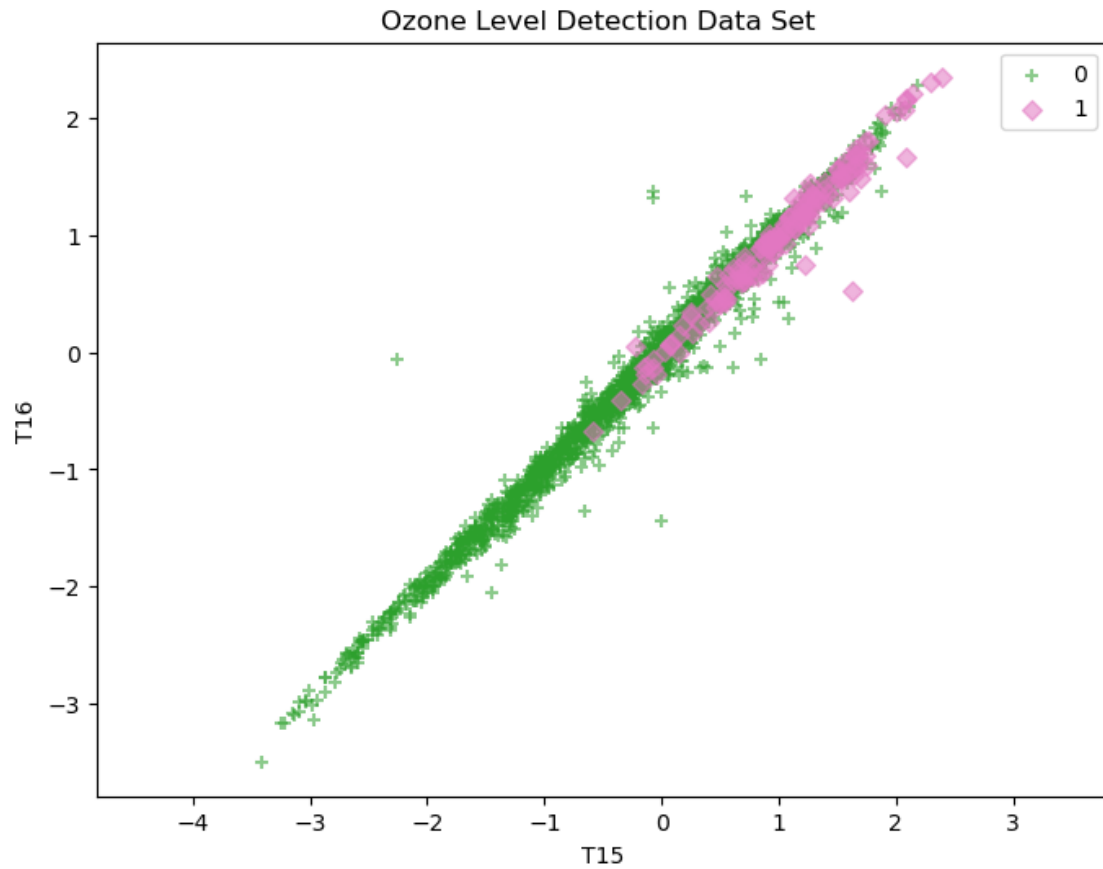
6.

```
[21]: y_int = y.astype(int)
```

```
[22]: plt.figure(figsize=(8, 6))
colors = ['tab:green', 'tab:pink']
markers = ['+', 'D']

for clr in y_int.unique():
    xx = data2[42].loc[y==clr]
    yy = data2[43].loc[y==clr]
    plt.scatter(xx, yy, c = colors[clr], label=clr,
                marker=markers[clr], alpha=0.55)
plt.axis('equal')
plt.title('Ozone Level Detection Data Set ')
plt.xlabel('T15')
plt.ylabel('T16')
plt.legend()
```

```
[22]: <matplotlib.legend.Legend at 0x2283d3b0610>
```



7. , 70% 30%. 3:
- :
- - 
  - (degree=2)

```
[23]: #
def train_test_split(X, y, test_ratio=0.2, seed=None):
    """returns X_train, X_test, y_train, y_test"""
    assert X.shape[0] == y.shape[0], \
        "the size of X must be equal to the size of y"
    assert 0.0 <= test_ratio <= 1.0, \
        "test_ratio must be valid"

    if seed:
        np.random.seed(seed)
```



```

shuffled_indexes = np.random.permutation(len(X))

test_size = int(len(X) * test_ratio)
test_indexes = shuffled_indexes[:test_size]
train_indexes = shuffled_indexes[test_size:]

X_train = X[train_indexes]
y_train = y[train_indexes]

X_test = X[test_indexes]
y_test = y[test_indexes]

return X_train, X_test, y_train, y_test

```

```

[24]: X_train, X_test, y_train, y_test = train_test_split(data2.to_numpy(), y, 0.3)

X_train.shape, y_train.shape, X_test.shape, y_test.shape

```

```

[24]: ((1774, 2), (1774,), (760, 2), (760,))

```

```

[25]: #
from sklearn.naive_bayes import GaussianNB

nbc = GaussianNB()
nbc.fit(X_train, y_train)
y_pred_nbc = nbc.predict(X_test)

```

```

[26]: from sklearn.linear_model import LogisticRegression

lg = LogisticRegression()
lg.fit(X_train, y_train)
y_pred_jg = lg.predict(X_test)

```

```

[27]: from sklearn.preprocessing import PolynomialFeatures

poly = PolynomialFeatures(degree=2)
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)

#
lgp = LogisticRegression()

#
lgp.fit(X_train_poly, y_train)

#
y_pred_lpg = lgp.predict(X_test_poly)

```

8.

. 6.

```
[28]: data2_np = data2.to_numpy()
```

```
[29]: import matplotlib.patches as mpatches
      from matplotlib.colors import ListedColormap
```

```
[31]: h = 0.01
      x_min, x_max = data2_np[:, 0].min() - 1, data2_np[:, 0].max() + 1
      y_min, y_max = data2_np[:, 1].min() - 1, data2_np[:, 1].max() + 1
      xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
      Z = nbc.predict(np.c_[xx.ravel(), yy.ravel()])
      Z = Z.reshape(xx.shape)

      cmap_decision = ListedColormap(['#C71585', '#00FFFF'])

      #
      plt.figure(figsize=(8, 6))
      plt.imshow(Z, interpolation="nearest", extent=(xx.min(), xx.max(), yy.min(), yy.
      ↪max()), aspect="auto", origin="lower", cmap=cmap_decision, alpha=0.3)

      colors = ['tab:green', 'tab:pink']
      markers = ['+', 'D']

      for clr in y_int.unique():
          xx = data2[42].loc[y==clr]
          yy = data2[43].loc[y==clr]
          plt.scatter(xx, yy, c = colors[clr], label=clr,
                      marker=markers[clr], alpha=0.55)


      legend_class0 = mpatches.Patch(color='#FF69B4', label='Class 0')
      legend_class1 = mpatches.Patch(color='#AFEEEE', label='Class 1')

      plt.legend(handles=[legend_class0, legend_class1], loc='lower right')

      plt.title("                                Ozone level detection")
      plt.xlim(x_min, x_max)
      plt.ylim(y_min, y_max)
      plt.xlabel('T15') # X
      plt.ylabel('T16') # Y
      plt.xticks(())
      plt.yticks(())

      plt.show()
```

Границы принятия решений для наивного байесовского классификатора Ozone level detection



```
[34]: h = 0.01
x_min, x_max = data2_np[:, 0].min() - 1, data2_np[:, 0].max() + 1
y_min, y_max = data2_np[:, 1].min() - 1, data2_np[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
Z = lg.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

cmap_decision = ListedColormap(['#C71585', '#00FFFF'])

#
plt.figure(figsize=(8, 6))
plt.imshow(Z, interpolation="nearest", extent=(xx.min(), xx.max(), yy.min(), yy.
    ↪max()), aspect="auto", origin="lower", cmap=cmap_decision, alpha=0.3)

colors = ['tab:green', 'tab:pink']
markers = ['+', 'D']

for clr in y_int.unique():
    xx = data2[42].loc[y==clr]
    yy = data2[43].loc[y==clr]
    plt.scatter(xx, yy, c = colors[clr], label=clr,
                marker=markers[clr], alpha=0.55)

legend_class0 = mpatches.Patch(color='#FF69B4', label='Class 0')
legend_class1 = mpatches.Patch(color='#AFEEEE', label='Class 1')
```

```

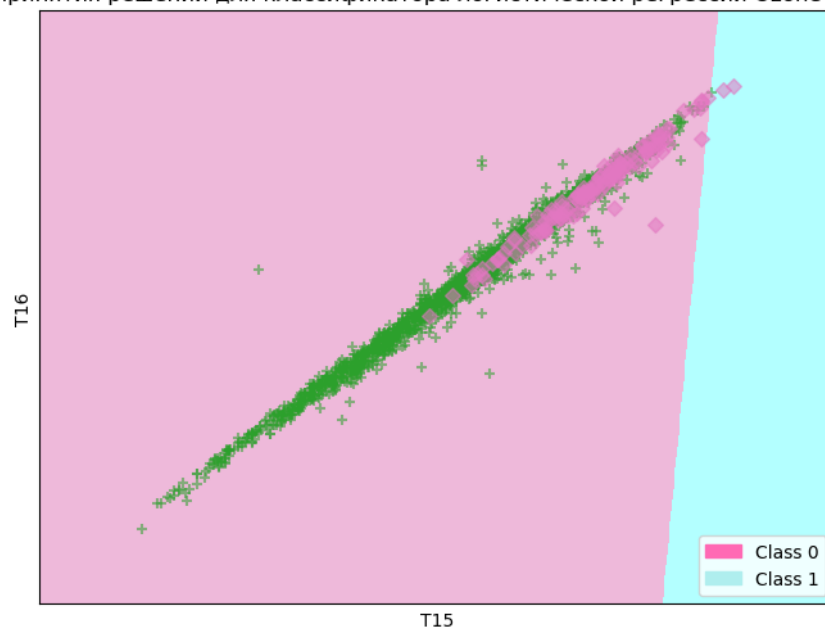
plt.legend(handles=[legend_class0, legend_class1], loc='lower right')

plt.title("
                                Ozone level detection")
plt.xlim(x_min, x_max)
plt.ylim(y_min, y_max)
plt.xlabel('T15')
plt.ylabel('T16')
plt.xticks(())
plt.yticks(())

plt.show()

```

Границы принятия решений для классификатора логистической регрессии Ozone level detection



```

[35]: x_min, x_max = data2_np[:, 0].min() - 1, data2_np[:, 0].max() + 1
      y_min, y_max = data2_np[:, 1].min() - 1, data2_np[:, 1].max() + 1
      xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                           np.arange(y_min, y_max, 0.01))

      #
      Z = lgp.predict(poly.transform(np.c_[xx.ravel(), yy.ravel()]))
      Z = Z.reshape(xx.shape)


      #
      plt.figure(figsize=(10, 6))
      plt.contourf(xx, yy, Z, alpha=0.8, cmap=plt.cm.RdYlBu)

```

```
#
scatter = plt.scatter(data2_np[:, 0], data2_np[:, 1], c=y, edgecolors='k',
    marker='o', cmap=plt.cm.RdYlBu)

#
plt.title("ozone level detection")
plt.xlabel("T15")
plt.ylabel("T16")
plt.legend(*scatter.legend_elements(), title=" ")
plt.show()
```

Граница принятия решения логистической регрессии с полиномиальными признаками ozone level detection



9. , 3: - ROC-

```
[36]: from sklearn.metrics import roc_curve, auc
y_prob_nbc = nbc.predict_proba(X_test)[: , 1]
y_prob_jg = lg.predict_proba(X_test)[: , 1]
y_prob_lpg = lgp.predict_proba(X_test_poly)[: , 1]

# ROC-
fpr_nbc, tpr_nbc, _ = roc_curve(y_test, y_prob_nbc)
fpr_jg, tpr_jg, _ = roc_curve(y_test, y_prob_jg)
fpr_lpg, tpr_lpg, _ = roc_curve(y_test, y_prob_lpg)

# AUC ( )
```

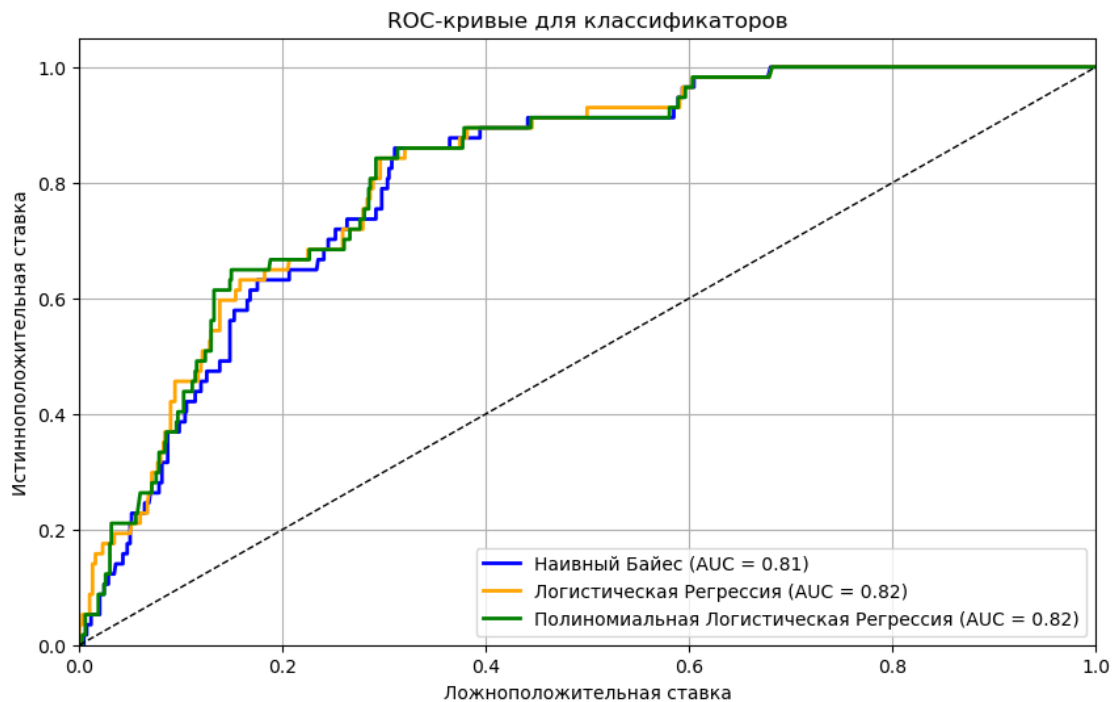
```

roc_auc_nbc = auc(fpr_nbc, tpr_nbc)
roc_auc_jg = auc(fpr_jg, tpr_jg)
roc_auc_lpg = auc(fpr_lpg, tpr_lpg)

# ROC-
plt.figure(figsize=(10, 6))
plt.plot(fpr_nbc, tpr_nbc, color='blue', lw=2, label='          (AUC = {:.2f})'.
        ↪format(roc_auc_nbc))
plt.plot(fpr_jg, tpr_jg, color='orange', lw=2, label='          (AUC = {:.
        ↪2f})'.format(roc_auc_jg))
plt.plot(fpr_lpg, tpr_lpg, color='green', lw=2, label='          □
        ↪(AUC = {:.2f})'.format(roc_auc_lpg))

#
plt.plot([0, 1], [0, 1], color='black', lw=1, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('          ')
plt.ylabel('          ')
plt.title('ROC-          ')
plt.legend(loc='lower right')
plt.grid()
plt.show()

```



10.

,

. 9.

```
[37]: # AUC
best_auc = max(roc_auc_nbc, roc_auc_jg, roc_auc_lpg)

if best_auc == roc_auc_nbc:
    best_model = " "
elif best_auc == roc_auc_jg:
    best_model = " "
else:
    best_model = " "

print(f" : AUC: {best_model}")
```

: AUC: